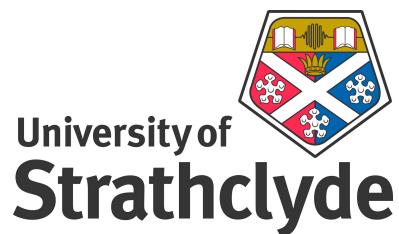


An Introduction to IEC 61970-301 & 61968-11: The Common Information Model



Dr Alan W. McMorran
Institute for Energy and Environment
Department of Electronic and Electrical Engineering
University of Strathclyde
Glasgow, UK

January 2007

The copyright of this publication belongs to the author under the terms of the United Kingdom Copyright Acts. Due acknowledgement must always be made of the use of any material contained in, or derived from, this publication.

1 Background

1.1 Introduction

Since deregulation, both in the UK and internationally, there has been an increasing need for power companies to exchange data on a regular basis. This is to ensure the reliable operation of the interconnected power networks owned and operated by a number of different utilities. Power companies use a variety of different formats to store their data, whether it be asset and work scheduling information in a proprietary internal schema within a database, topological power system network data within a control system, or static files used by simulation software.

While much of this data is only required within a company, there is often a need to exchange the data both internally between different applications and externally with other companies. The large number of proprietary formats used by these applications requires a myriad of translators to import and export the data between multiple systems. This exponential growth in complexity when integrating increasing numbers of applications and exchanging between multiple companies has driven the requirement for a common format that covers all the areas of data exchange in the power electrical domain.

The IEC standard 61970-301 [1] is a semantic model that describes the components of a power system at an electrical level and the relationships between each component. The IEC 61968-11 [2] extends this model to cover the other aspects of power system software data exchange such as asset tracking, work scheduling and customer billing. These two standards, 61970-301 and 61968-11 are collectively known as the Common Information Model (CIM) for power systems and currently have two primary uses: to facilitate the exchange of power system network data between companies; and to allow the exchange of data between applications within a company.

The development of the CIM has primarily taken place in North America, where the North American Electric Reliability Council (NERC) has adopted the CIM as the format for exchanging network data between transmission companies. The majority of the application integration activities have similarly taken place within North American utilities.

This paper describes the pitfalls of the traditional methods of storing power system data, then introduces the concepts behind class modelling and how this approach is used to define a power system in the IEC 61970 Common Information Model (CIM) standard. The use of the Extensible Markup Language (XML) to encapsulate this data for the exchange of both full power system models and inter-application messages is then described.

1.2 Power System Data Formats

Since the advent of the modern digital computer, power system engineers have utilised the capabilities of this tool in a variety of areas, whether it be performing complex analysis calculations on a power system or to control its operations in real-time. All of these applications require the operator to digitally store and exchange data about the system.

Large-scale Energy Management Systems (EMS) and asset-management systems use database schemas for defining the structure of the data storage data, often custom-written to reflect the operator's specific requirement. Offline applications for performing load-flow and fault-level analysis simulations use application-specific file formats that represent the data required by each application.

In modern utilities' IT infrastructures, large-scale applications such as the EMS and asset-management system communicate with each other, generally using a vendor's own custom format based on the internal database schema. In the past this often required the user to purchase each piece of enterprise-level software from the same vendor to ensure compatibility when integrating them.

The deregulation of the power industry, however, has resulted in multiple utilities, running software from a number of different vendors, having to exchange large data sets on a regular basis. The use of proprietary, custom formats complicates this exchange, requiring complex translation between each of the custom formats.

Similarly, offline applications traditionally use a rigid, proprietary format containing only the data required by that particular version of the application. When subsequent versions of the program require additional details the file format is changed, resulting in multiple formats for a single application. Of course, such a scenario is not limited to power system applications. Changing the file format for each new software version is common practice within the software industry but usually only causes minor irritation since each new version of a vendor's software contains import facilities to convert previous versions of the file format into the new format.

Problems occur when companies need to exchange data between software applications from different vendors, and / or have multiple versions of the same software running within their company. Such a scenario requires a company to either:

1. Maintain multiple copies of the same data in multiple formats
2. Store the data in a format compatible with every piece of software, requiring the removal of application-specific data and a subsequent loss in precision
3. Store the data in a single, highly-detailed format and create software to translate from this highly-detailed format to the desired application file formats
4. Use a highly detailed format that is compatible with every application and whose standard format contains the basic data required to represent the power system while simultaneously allowing additional, detailed, application-specific data to be contained without invalidating the format.

The third option requires additional software engineering on the part of the company to create translation tools, but requires them to maintain only a single format containing all the data required. The fourth option represents the ideal solution, allowing a company to maintain a single, highly detailed format that is compatible with any of their software.

This option does, however, requires three things:

- A highly detailed model to describe the power system
- A file format capable of storing extended data without affecting the core data
- Power system software vendors and utilities to either adopt and embrace this data model and format either for economic or regulatory reasons

The Common Information Model (CIM) for Power Systems has the potential to meet the first requirement of the above list while the eXtensible Markup Language (XML),

combined with the Resource Description Framework (RDF) offers a means of fulfilling the second requirement. The remaining requirement can be considered more of a commercial challenge than a technical one. Universal acceptance of this format requires both utilities and vendors to acknowledge the benefits of adopting the standard. At present, all of the major power system application vendors are active participants in the CIM Interoperability tests and the popularity of the format is spreading.

This paper will provide some background on the CIM and the CIM RDF XML format. To understand the structure of the CIM, however, it is important to have an understanding of class hierarchies within the object-oriented software paradigm and the benefits of using such an approach to model the components of a power system. The following sections will provide some general background on class hierarchies, followed by more detailed background information on the CIM. Finally, it will be shown how this data can be represented in the RDF XML format.

2 Class Hierarchies and UML Class Diagrams

When building any system to represent data, whether it be a software architecture or a database schema, the design of the system will define how extensible and scalable the system is, and ultimately, whether it succeeds or fails at its given task. This paper provides an introduction to the concept of Class Hierarchies and how they are used in system design, along with the Unified Modelling Language (UML).

Within a system, a class represent a specific type of object being modelled. A class hierarchy is an abstract model of a system defining every type of component within a system as a separate class. A class hierarchy should reflect the real-world structure of the system.

While a full description of UML is outwith the scope of this thesis, UML class diagrams provide a useful means of visually representing object hierarchies. This section will provide a simple case study to show how a class hierarchy representing a small segment of a University system can be constructed independently of the final platform on which the design will be utilised.

2.1 Classes

Each class can have its own internal attributes and relationships with other classes. Each class can be instantiated into any number of separate instances, known as objects, each containing the same number and type of attributes and relationships, but with their own internal values.

Person
Name
Gender

Figure 2.1 The Person Class

A simple example of a class is that of a Person as shown in Figure 2.1. The Person class contains two basic attributes: *Name* and *Gender*. If the system being created were to represent every person in the University, it would require only this single class since every person within the University can be represented at the most basic level by the attributes defined in Person.

For a University containing 10,000 students and staff, the system would create 10,000 separate instances of the Person class, each containing a value for Name and Gender independent of the other 9,999 instances (although not necessarily unique).

If the system is required to store more information than just a person's Name and Gender, and differentiate between staff, students and the different types of each, then the class diagram becomes more complex.

2.2 Inheritance (Generalisation)

Inheritance (also known as Generalisation) defines a class as being a sub-class of another class. As a sub-class, it inherits all the attributes of its parent, but can also contain its own attributes.

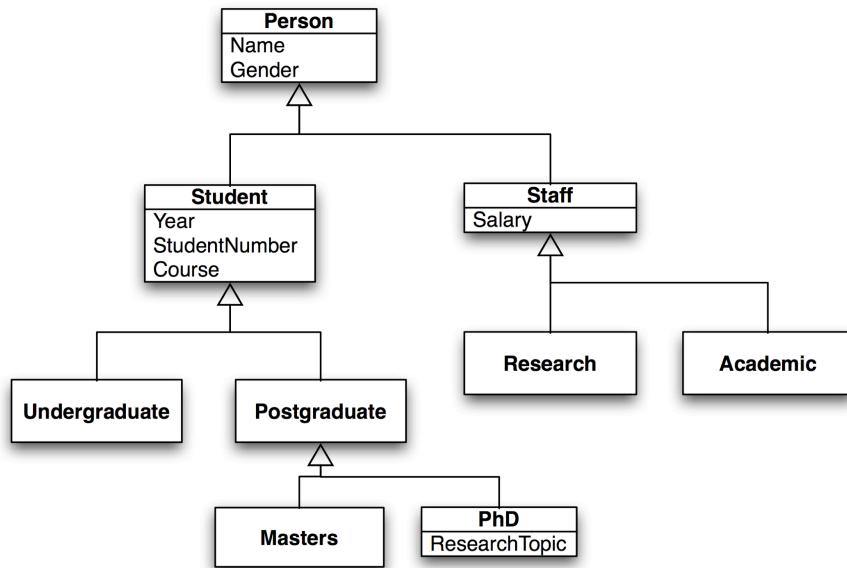


Figure 2.2 Class Hierarchy of people at a University

Figure 2.2 provides a class hierarchy to represent some of the different types of people that exist within a University system. This diagram, as with all subsequent class diagrams uses standard UML symbology. Student and Staff are both sub-classes of Person. A Student is still a person and still has a Name and Gender, but has additional attributes to denote the year they are in, their student number and the course they are studying. Similarly, if someone is Staff they are still a Person, but have gained a new attribute to indicate their salary.

The Student class itself has two sub-classes, Undergraduate and Postgraduate, both inheriting all the attributes of Student (and in turn, of Person), but independent of each other. The Postgraduate class also has two subclasses, Masters and PhD. A PhD is a Postgraduate and a Student and a Person, with all of their attributes, but with the addition of its own **ResearchTopic** attribute.

The Staff class also has two sub-classes, Research and Academic, both of which retain all the attributes of Staff and Person.

So while a PhD is a Postgraduate, a Student and a Person, not all Students are PhDs. Similarly, an Academic is Staff and a Person, but not every member of Staff is an Academic.

2.3 Association

Section 2.2 has shown how a class hierarchy can be formed to describe sub-classes of Person, but other than the inheritance there are no other relationships defined between classes the classes in Figure 2.2.

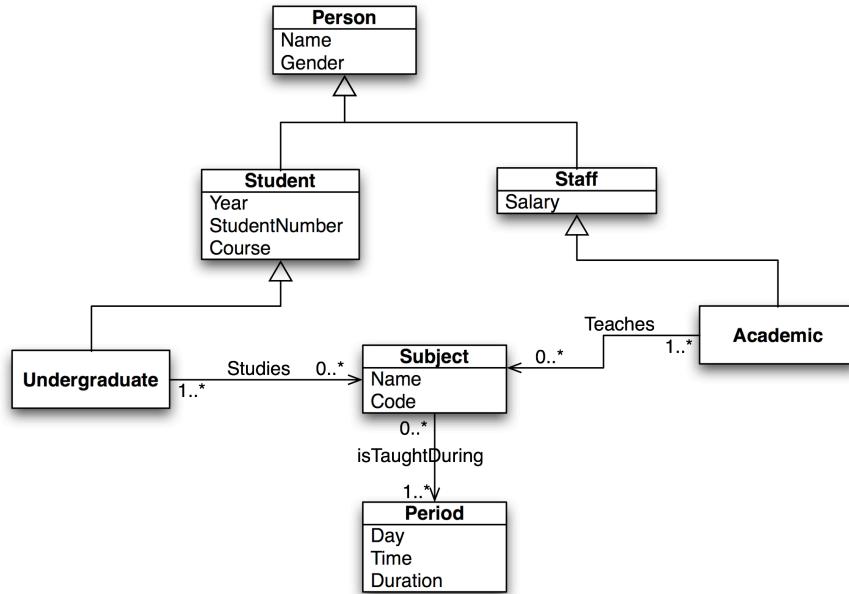


Figure 2.3 Class hierarchy of students, staff and subjects

In Figure 2.3 we have introduced two additional classes: Subject and Period. Neither of these classes are a type of Person, and as such do not inherit from the Person class. The Subject class does however, have a relationship with the Undergraduate and Academic classes.

An Undergraduate can study a number of subjects and an Academic can teach a number of subjects. These relationships are shown on the diagram as associations between the classes.

For the Undergraduate-Subject association, the role is given as “Studies” while the location of the arrowhead indicates that it is the Undergraduate who *Studies* the Subject (if the arrow were reversed, it would mean that the Subject studied the Undergraduate).

At each end of the association link is the multiplicity. For the Undergraduate-Subject association, these indicate that a Subject must have from 1 to many (**1..***) Undergraduates, but an Undergraduate can have from 0 to many (**0..***) subjects (since it is possible that some Undergraduates may not be studying any subjects due to industrial placements, sabbaticals etc. It is assumed that a subject will not exist in the system if no students have chosen to study it).

Similarly, a Subject will have from 1 to many Academics who teach that Subject, but an Academic may teach from 0 to many Subjects (since not all Academics have to teach).

The other additional class, Period, with Day, Time and Duration attributes, represents a particular timetable period and as such a Subject has an association with a Period. As with the other associations, the Subject-Period association has a role, *isTaughtDuring* and

multiplicities which indicate that a Subject will be taught during 1 to many periods and that a Period will have from 0 to many Subjects taught during its time-period.

This demonstrates how classes relate to each other on a very basic level and how UML Class Diagrams provide a means of graphically displaying these relationships.

2.4 Aggregation

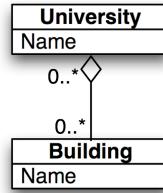


Figure 2.4 Class Hierarchy of a University and Building

The Aggregation relationship defines a special kind of association between classes, indicating that one is a container class for the other. In the example shown in Figure 2.4, two new classes University and Building have been introduced, each very simple classes containing only a single attribute to denote their name. The multiplicity on the diagram operates in the same manner as to that of the Association, indicating that a Building can be part of 0 or more Universities (we are assuming that some Universities operate joint schools and not every building within the system will necessarily be part of a University). The second multiplicity indicates that a University can contain 0 or more buildings (0 if it operates solely by remote learning for example).

Unlike a simple Association relationship the line denoting a relationship on the diagram contains a diamond instead of an arrowhead. This indicates that the two classes have an Aggregation relationship. This can be thought of as “The University is made up of 0 or more Buildings”, indicating that the relationship is stronger than a simple association. The clear diamond, however, indicates that the two are not completely inter-dependent, and that if the University were destroyed the buildings would still exist (assuming the destruction was not a literal demolition but instead indicated that the University had ceased to exist).

2.5 Composition

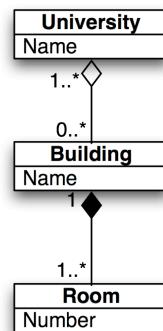


Figure 2.5 Class Hierarchy of a University, Building and Room

Composition is a specialised form of Aggregation where the “contained” object is a fundamental part of the “container” object, and that if the “container” is destroyed, all the objects that are related to it with a composition are similarly destroyed. An example of this is shown in Figure 2.5 between the new Room class and the Building class.

The line here has a solid diamond, indicating that the relationship is a composition. The multiplicity states that a building will have 1 or more rooms (since even an empty building can be thought of as one giant room) and that a room will be contained within 1 building only. This reflects the real world makeup of rooms and buildings. If a building is destroyed then the rooms within it are also destroyed.

Any system that implements this design will know that if a Building object is destroyed, any Room objects that are contained within that particular instance will also be destroyed.

2.6 Summary

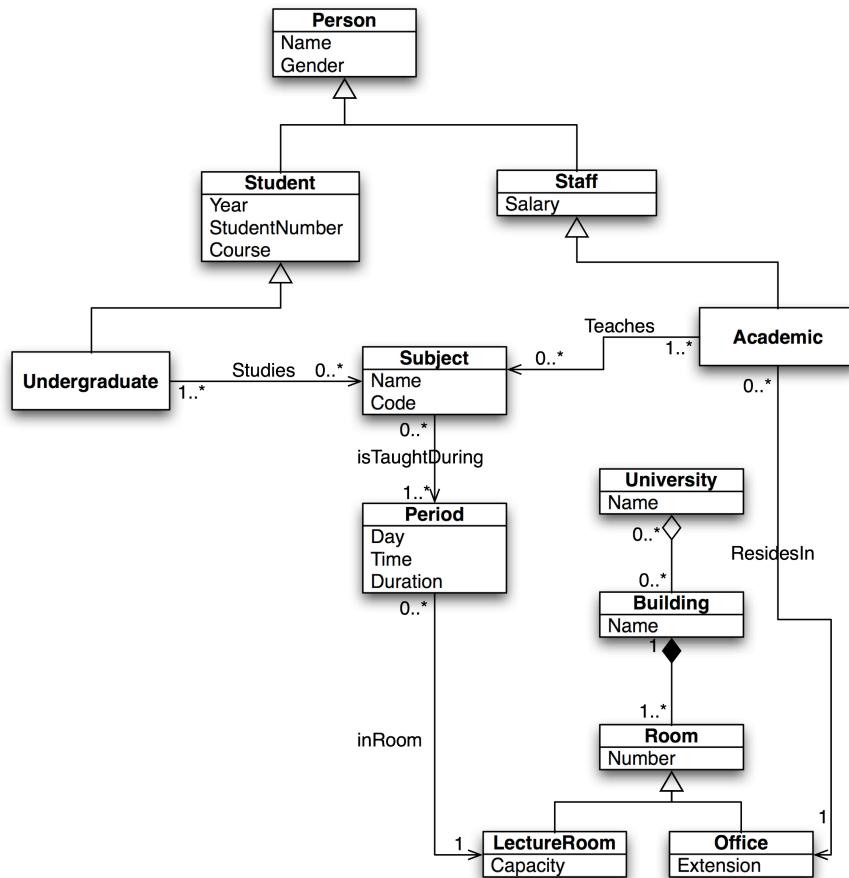


Figure 2.6 Class diagram showing some of previous classes and their relationships

The previous sections should have provided you with a basic understanding of what a class hierarchy is and how this can be represented on a class diagram. Figure 2.6 shows some of the classes from the previous sections (along with two extra sub-classes of Room), and how the separate diagrams in Figure 2.3, Figure 2.4 and Figure 2.5 all relate to each other. It should be clear that the system could be extended further to incorporate more details about the University system such as timetables for students

and staff or the computing facilities available in each room. Both of these examples would make use of the existing classes by association along with the introduction of new classes.

These fundamentals of the class system are essential in the understanding of the CIM as described in the following sections.

3 The Common Information Model for Power Systems

This section provides some history of the CIM and how it represents the common components within a power system. Examples are given to show where a common power system component fits into the class hierarchy, how connectivity is represented using CIM classes and how an example circuit, shown as a line diagram, can be converted to CIM objects. Finally, each of the packages in the IEC 61970-301 standard is summarised.

3.1 History

Exchanging power systems data between utility companies is always problematic when proprietary formats are used. In the past, a company would traditionally use a single software system, whether it is a custom in-house solution, or purchased from a large software company, and there would be a single proprietary data standard and format used. With the deregulation of the power industry both in the UK and abroad, there is now a greater need to be able to share such power system data between companies. The increase in choice provided by the number of power system software vendors, and the different software packages and architectures available add to the challenge of data exchange. These issues point to a requirement for a single, open standard for describing power system data and to aid the interoperability between software packages and exchange of information both within one company and between companies.

The Common Information Model (CIM)[1][2] is an open standard for representing power system components developed by the Electric Power Research Institute (EPRI) in North America. The standard was developed as part of the IEC TC57 WG13 on developing a Control Centre Application Programming Interface (CCAPI) to provide a common model for describing the components in power systems for use in a common Energy Management System (EMS) Application Programming Interface (API). The format has been adopted by the major EMS vendors to allow the exchange of data between their applications, independent of their internal software architecture or operating platform.

The data model itself is language-independent, defining the components of a power system as classes along with the relationships between these classes: inheritance, association and aggregation; and the parameters within each class are also defined. This provides the foundation for a generic model to represent all aspects of a power system, independent of any particular proprietary data standard or format. This simplifies the interoperability between software applications, since there need only exist a translator to convert to and from the CIM based data, where previously there would have been the need for translators to convert to and from every other third party company's proprietary format.

For an engineer the format of the Common Information Model (CIM) may at first appear confusing compared with a flat file format. This paper will explain how the CIM was created using a class structure to describe components of a power system network; the advantages of this approach; and how a power system network model can be translated into a number of CIM objects.

3.2 CIM Class Structure

The CIM hierarchy currently has no official common super-class (i.e. a class from which every component inherits). The majority of CIM classes, however, inherit from the IdentifiedObject class so for this section it can be considered the base class for the hierarchy.

3.2.1 Example: The Breaker Class

A simple example will be used to explain why it is advantageous to use a class structure for defining components instead of simply specifying attributes for every different type of component in the CIM as an independent entry.

A Breaker is one of the most common components in a power system described as a “mechanical switching device capable of making, carrying and breaking currents under normal circuit conditions and also making, carrying for a specified time, and breaking current under specified abnormal circuit conditions”[1]. To understand how this fits into the CIM class hierarchy the Breaker can be thought of at different levels of abstraction.

At the most detailed level it is a Breaker, but since a breaker’s most basic functionality is the ability to be open or closed it can be described as a specialised type of switch. Within the power system a switch is part of the physical network that conducts electricity, and as such can be considered a type of conducting equipment. Since the power system may contain equipment that does not conduct electricity directly, conducting equipment can be considered a type of generic equipment. A piece of equipment can similarly be considered as a being resource within the power system.

A Breaker can therefore be considered to be a Power System Resource, a type of Equipment, a type of Conducting Equipment and a type of Switch. This corresponds to a class inheritance structure shown in Figure 3.1 below.

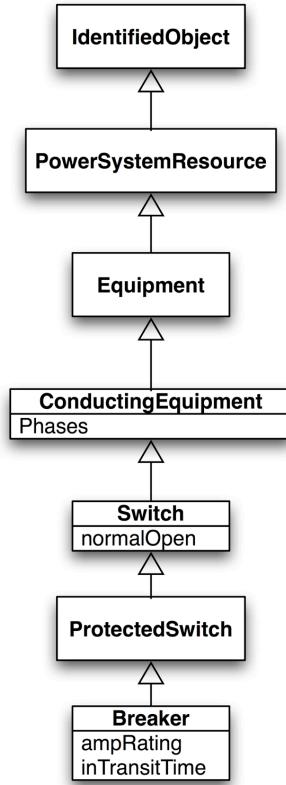


Figure 3.1 Breaker Class Inheritance Hierarchy

The Naming class is the root class for this particular branch of the CIM class hierarchy and other CIM classes in the Breaker hierarchy are:

- PowerSystemResource, used to describe any resource within the power system, whether it be a physical piece of equipment such as a Switch or an organisational entity such as a SubControlArea.
- Equipment, which refers to any piece of the power system that is a physical device, whether it be electrical or mechanical.
- ConductingEquipment, used to define types of Equipment that are designed to carry current or that are conductively connected to the network and contains an attribute to denote the phases (A,B,C,N or any combination of each).
- Switch, a generic class for any piece of conducting equipment that operates as a switch in the network and hence has an attribute to define whether the switch is normally open or closed.
- ProtectedSwitch, a switch that can be operated by protection equipment
- Breaker, a specific sub type of ProtectedSwitch, with additional attributes to define the current rating and transit time.

As with the University system example in Section 2, all subclasses inherit the attributes from their parent class, and as such a Breaker will contain a **normalOpen**, from the **Switch** class, and **phases** attribute, from the **ConductingEquipment** class, as well as its own native attributes.

3.2.2 Subclasses of Switch

As well as Breaker, the CIM standard contains multiple subclasses of Switch, including Jumper, Fuse, Disconnector, LoadBreakSwitch and GroundDisconnector.

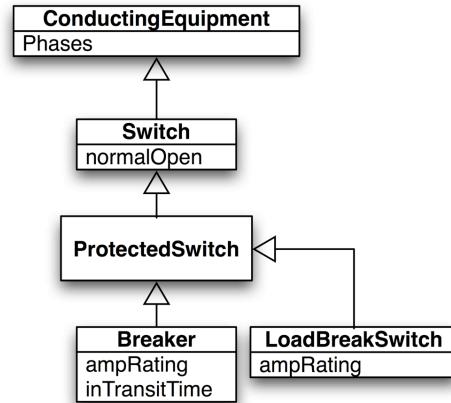


Figure 3.2 Switch class with Breaker and LoadBreakSwitch subclasses

Figure 3.2 shows an example of how the LoadBreakSwitch class, a subclass of ProtectedSwitch fits into the class hierarchy. Both Breaker and LoadBreakSwitch inherit from ProtectedSwitch which in turns inherits from Switch, so they both contain a *normalOpen* attribute whilst maintaining their own, internal attributes.

As well as dealing with them as their native class, the system can treat a Breaker or LoadBreakSwitch component as being a ProtectedSwitch, Switch, a piece of Conducting Equipment, a piece of Equipment, a Power System Resource or just an IdentifiedObject.

For example:

If a piece of software is performing a topological analysis on a power system network then it will need to know whether a switch is open or closed to determine the status of the network. The software does not need to know whether the Switch is a Breaker, a LoadBreakSwitch or any other subtype of Switch since the attribute it is concerned with, *normalOpen*, exists in all the classes that inherit from Switch. As the software traverses the network model, if the component it reaches is of the class Switch or any of its subclasses it extracts the value of *normalOpen* and proceeds accordingly.

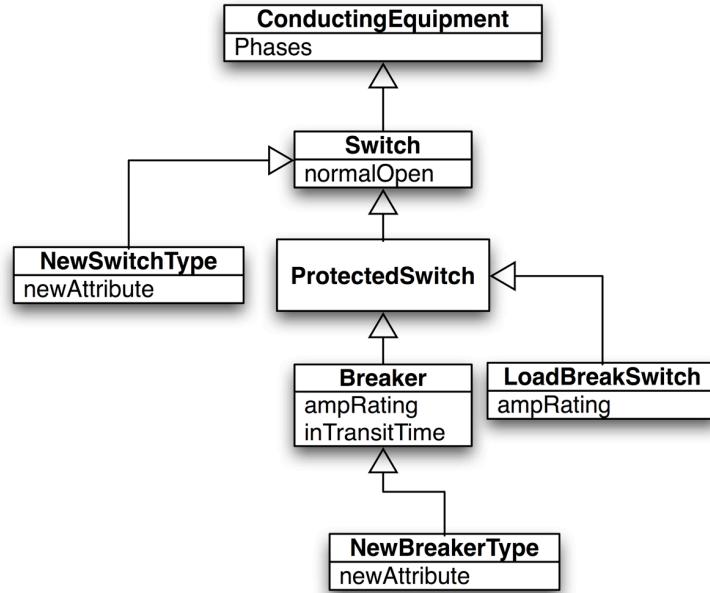


Figure 3.3 Switch Class diagram with new subclasses of Switch and Breaker

If a new type of Switch, **NewSwitchType** is added to the standard at a later date as shown in Figure 3.3, assuming the original **Switch** class is not modified, then the software will still be able to treat **NewSwitchType** as if it were a **Switch** when performing its analysis. Even though the class did not exist when the software was originally written it is looking for any components that are of a class that inherits from **Switch**.

Similarly, if a new subclass of Breaker, **NewBreakerType**, is added (as shown in Figure 3.3), it is still a type of **Switch** (since its parent class, **Breaker** is a subclass of **ProtectedSwitch** which itself is a subclass of **Switch**) and can be treated as **Switch**, **ProtectedSwitch** or a **Breaker** by the software.

As has been shown, this use of an inheritance hierarchy to define components allows classes within the system to be defined as specialised subclasses of a general parent class until the desired level of detail has been reached, from the generic **PowerSystemResource** right down to the **Breaker** or **LoadBreakSwitch** class.

This use of a class hierarchy also allows extensions to be made to the standard by extending the existing classes instead of introducing completely new, independent entries. This approach, as shown, can allow existing software applications to interpret the new data, albeit at a higher level of abstraction, without necessarily requiring extensive modification.

3.2.3 Defining Component Interconnections

When defining how components within a power system network join together, rather than define direct connection between components, the CIM uses Terminals and Connectivity Nodes.

To understand why this approach is taken consider the very simple, circuit shown in Figure 3.4 below.

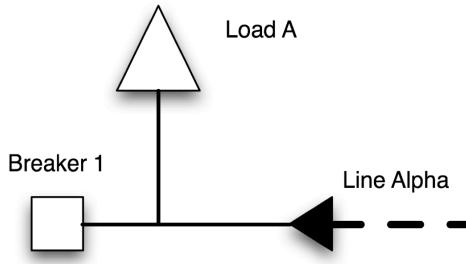


Figure 3.4 Connectivity Example circuit

This circuit, containing a Breaker, Load and Line, would require three CIM Objects to represent the pieces of physical conducting equipment: An Energy Consumer (to represent the load), a Breaker and an AC or DC Line Segment for the line.

The CIM does not model interconnections by associating each component with the other components it connects to, since having Breaker 1 contain associations to Load A and Line Alpha; Load A contain associations to Line Alpha and Breaker 1; and Line Alpha contain associations to Breaker 1 and Load A would result in the interconnections being defined as shown in Figure 3.5.

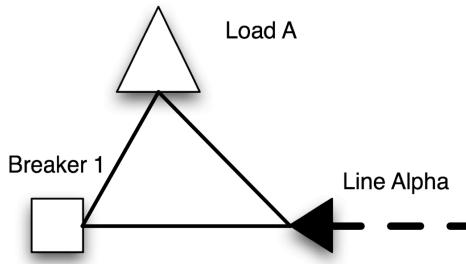


Figure 3.5 Connectivity Example circuit with direct associations

Instead, the CIM uses a Connectivity Node to connect equipment, so that should three or more pieces of equipment meet at a T or Star point, the connectivity is accurately represented as shown in Figure 3.6.

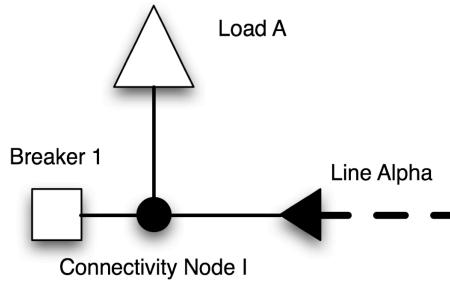


Figure 3.6 Connectivity Example circuit with Connectivity Node

In CIM, however, pieces of conducting equipment are not directly associated with Connectivity Nodes. A piece of conducting equipment will have one or more Terminals associated with it, and these Terminals in turn are associated with a single Connectivity Node.

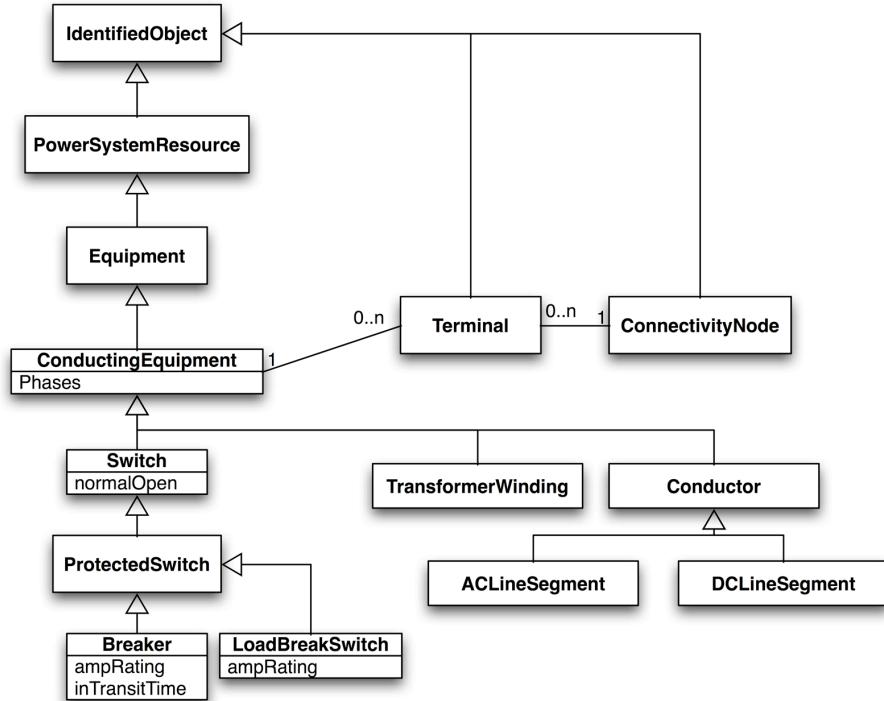


Figure 3.7 Conducting Equipment and Connectivity class diagram

The relationship between the Terminal, ConnectivityNode and ConductingEquipment classes is shown in Figure 3.7. Since only pieces of conducting equipment carry current on the network, the association to the Terminal class is from the ConductingEquipment class with a multiplicity of 0..n since a piece of conducting equipment can have zero or more connections to the network. The corresponding Terminal to Conducting Equipment relationship has a multiplicity of 1 since a Terminal can only ever be associated with one Connectivity Node. Since the Breaker class (via its Switch class parent), Energy Consumer and AC or DC Line Segment (via the Conductor class) all inherit from Conducting Equipment, they too inherit the association relationship with the Terminal class.

The connectivity relationship between the terminals, conducting equipment and connectivity nodes is illustrated in Figure 3.8a) below.

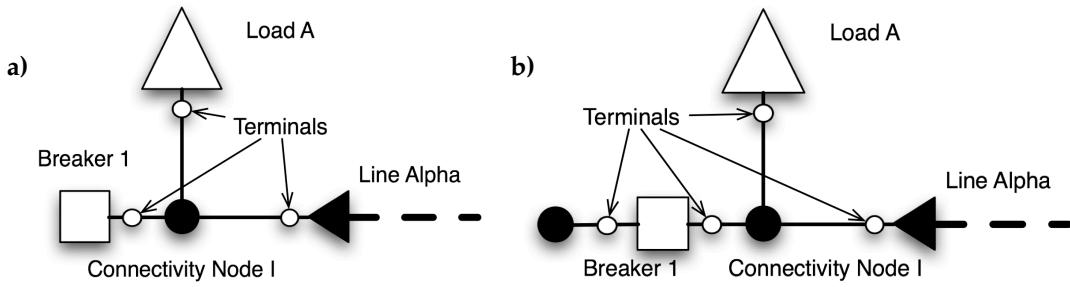


Figure 3.8 Connectivity Example circuit with Connectivity Node and Terminals

The inclusion of the Terminals may initially seem unnecessary, but as well as defining connectivity, Terminals are also used for defining points of connectivity-related measurement in the network such as power flows, currents and voltages.

The importance of allowing the measurement point to be defined so exactly can be shown in Figure 3.8b). In this diagram Breaker 1 has two Terminals associated with it to represent the two distinct network connection points it would have in a real-world power system network. If the Breaker is open then the measurement of voltage for the Breaker will be different at these two points where the Breaker connects to the network. This would result in an ambiguity if measurement were only defined as being on a particular component without specific information about which point of connection the measurement is to be made at.

3.3 Converting a Circuit to CIM Objects

The previous sections have described a small section of the class hierarchy for describing CIM components and shown how Terminals and Connectivity Nodes are used to define the interconnection of components within the network. This section will use a more complex example to show how voltage levels, current transformers, power transformers and generators are modelled by converting a standard line diagram into CIM objects.

3.3.1 Identifying the CIM Classes

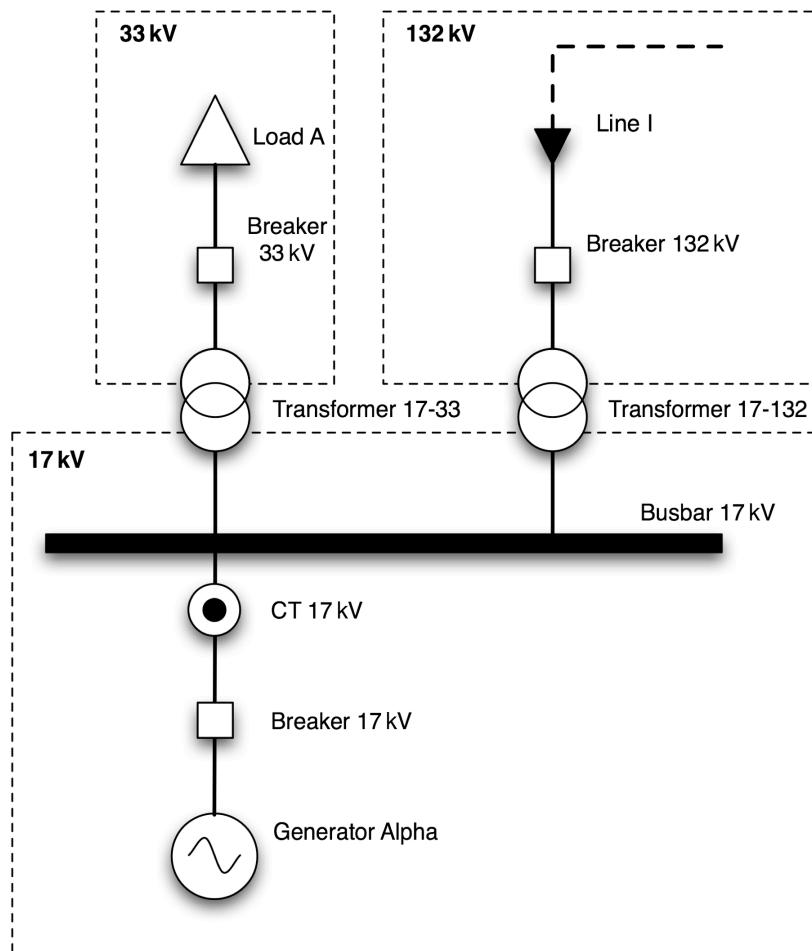


Figure 3.9 Example Circuit as a line diagram

The circuit shown in Figure 3.9 shows a circuit containing a single generating source, load, line and busbar. The circuit also contains two power transformers resulting in three distinct voltage levels of 17kV, 33kV and 132kV.

The load, line and breakers, as stated in Section 3.2.3 map to the CIM EnergyConsumer, ACLineSegment and Breaker classes respectively while the busbar similarly maps to the BusbarSection class. Generator Alpha will map to a single piece of conducting equipment, the SynchronousMachine, an “electromechanical device that operates synchronously within the network”[1]. When operating as a generator, the SynchronousMachine object must have an association with an instance of the GeneratingUnit class.

The GeneratingUnit class does not represent a piece of conducting equipment that physically connects to the network; instead it represents “a single or set of synchronous machines for converting mechanical power into alternating-current”[1]

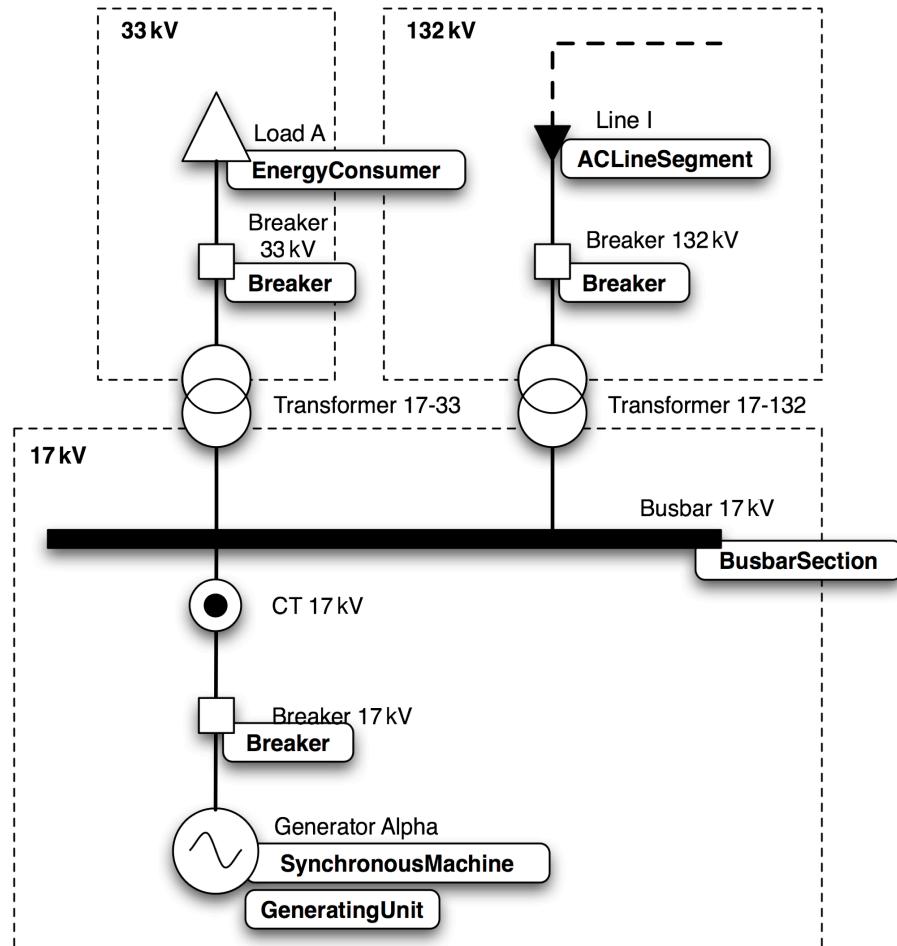


Figure 3.10 Example Circuit with partial CIM Class mappings

These mappings are shown in Figure 3.10, leaving only the two power transformers and current transformer to be mapped to CIM classes.

3.3.2 Representing Power Transformers as CIM Objects

A power transformer is not mapped to a single CIM class, instead it is split down into a number of components with a single PowerTransformer *container* class. Thus a two-winding power transformer becomes two TransformerWinding objects within a PowerTransformer container. If a tap changer is present to control one of the windings then an instance of the TapChanger class is associated with that particular winding while still being contained within the PowerTransformer instance. The UML class diagram for the classes that form a transformer is shown in Figure 3.11 below.

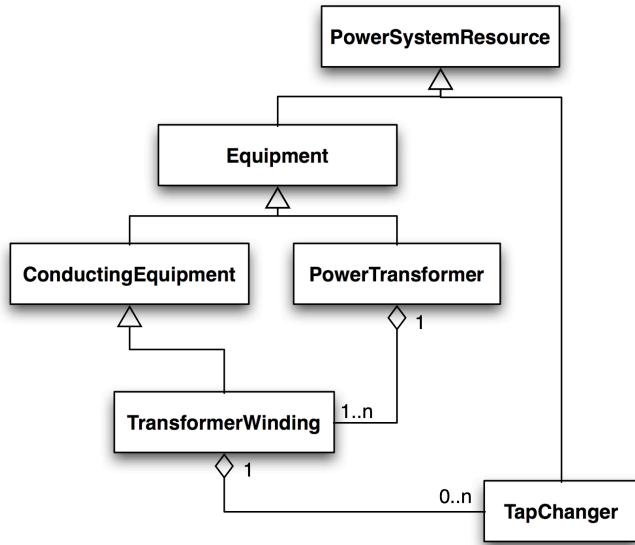


Figure 3.11 Transformer Class Diagram

Although a PowerTransformer is still a piece of Equipment in the system, it does not conduct electricity itself and thus does not inherit from ConductingEquipment but from its parent, Equipment. A TransformerWinding, however, does inherit from ConductingEquipment since it is physically connected to the network and does conduct electricity. The TapChanger is part of the TransformerWinding and as such cannot be considered to be a separate piece of equipment in its own right and inherits from PowerSystemResource.

The PowerTransformer and TransformerWinding classes have an aggregation relationship¹, meaning that a PowerTransformer is made up on 1 or more TransformerWindings which in turn can be made up of zero or more TapChangers.

When considering a physical transformer sitting in a substation the PowerTransformer container can be thought of as the shell of the transformer. The shell itself does not conduct any of the electricity in the network, but instead holds the windings of the transformer, the insulating material, magnetic core, and all the other components that make up the transformer.

¹ Although it could be argued that this relationship is composition rather than aggregation the CIM class structure contains no composition relationships. This is due to the flexible design of the standard, where a composition relationship would indicate a tighter relationship between classes than is necessary for a number of applications of the standard.

The connections from the transformer to the network are made with the windings themselves, a relationship that is mirrored in the CIM representation where it is the TransformerWinding class that inherits from ConductingEquipment.

Transformer 17-33

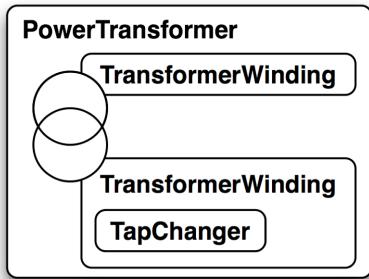


Figure 3.12 CIM Mappings for Transformer 17-33

Thus, Transformer 17-33 from Figure 3.9 can be represented as 4 CIM objects: two TransformerWindings, one TapChanger and one PowerTransformer as shown in Figure 3.12.

Similarly, a transformer with a tertiary or quartary winding can be represented as a single PowerTransformer containing three or four instances of the TransformerWinding class.

3.3.3 Representing a Current Transformer as a CIM Object

The current transformer CT 17kV does not map directly to a piece of conducting equipment in the CIM hierarchy as would be expected. The current transformer's purpose is to measure the current at its location in the network, and as such when modelling the network it is the measurement from that location that is modelled rather than the piece of equipment doing the measuring.

This involves creating an instance of the Measurement class to measure the current at a particular terminal. As described in Section 3.2.3, each piece of conducting equipment has one or more terminals to represent the points at which it connects to the network. By associating a Measurement object with a particular terminal and defining the measurement taken by that instance to be current then the Measurement object will reflect the role played by the current transformer.

3.3.4 Defining Containment

As well as having component interconnections defined using the ConductingEquipment-Terminal-ConnectivityNode associations, the CIM has an EquipmentContainer class that provides a means of grouping pieces of Equipment together to represent both electrical and non-electrical containment.

Voltage Levels

Pieces of conducting equipment do not have a voltage attribute to define the voltage as a specific value, instead they are associated with a VoltageLevel, a subclass of EquipmentContainer. Each instance of the VoltageLevel class itself has an associated

BaseVoltage object that contains a single attribute to define the nominal voltage of that particular group of components. A BaseVoltage instance may be associated with more than one VoltageLevel, since standard voltage levels (e.g. 33, 132, 275, 400kV) will exist throughout the network. Each VoltageLevel instance, however, contains only the interconnected pieces of equipment at the same voltage level. This is an example of using a subclass of EquipmentContainer to represent electrical containment.

Substations

The Substation class is a subclass of EquipmentContainer that can contain multiple VoltageLevels and is used to define a collection of equipment “through which electric energy in bulk is passed for the purposes of switching or modifying its characteristics”[1].

In the example network shown in Figure 3.9, the three different voltage levels identified by the dashed bounding boxes are mapped to three instances of the VoltageLevel and contained within a single SubStation instance. Each VoltageLevel object also has an associated BaseVoltage object with a nominal voltage of 17, 33 and 132kV.

The Substation class, being a subclass of EquipmentContainer can also contain other instances of Equipment, such as PowerTransformer, which, as previously explained, is itself a container, not a piece of conducting equipment. The Substation class is an example of a subclass of EquipmentContainer to represent non-electrical containment since it will contain pieces of equipment that are physically grouped, but not necessarily electrically connected.

Lines

The ACLineSegment, however, is not contained within a VoltageLevel, instead it is contained within an instance of the Line class. The Line class in CIM is used to define a “component part of a system extending between adjacent substations or from a substation to an adjacent interconnection point”[1]. A Line may contain multiple line segments of either the AC or DC variety, but does not itself represent a piece of physical conducting equipment.

Since a line segment is used to represent “a wire or combination of wires ... used to carry alternating [or direct] current between points in the power system”[1] it would be inaccurate to define it as being inside a specific voltage level within a substation. As such, the AC and DCLineSegment classes contain a direct association to the BaseVoltage class to define their nominal voltage level.

3.3.5 Equivalent CIM Representation

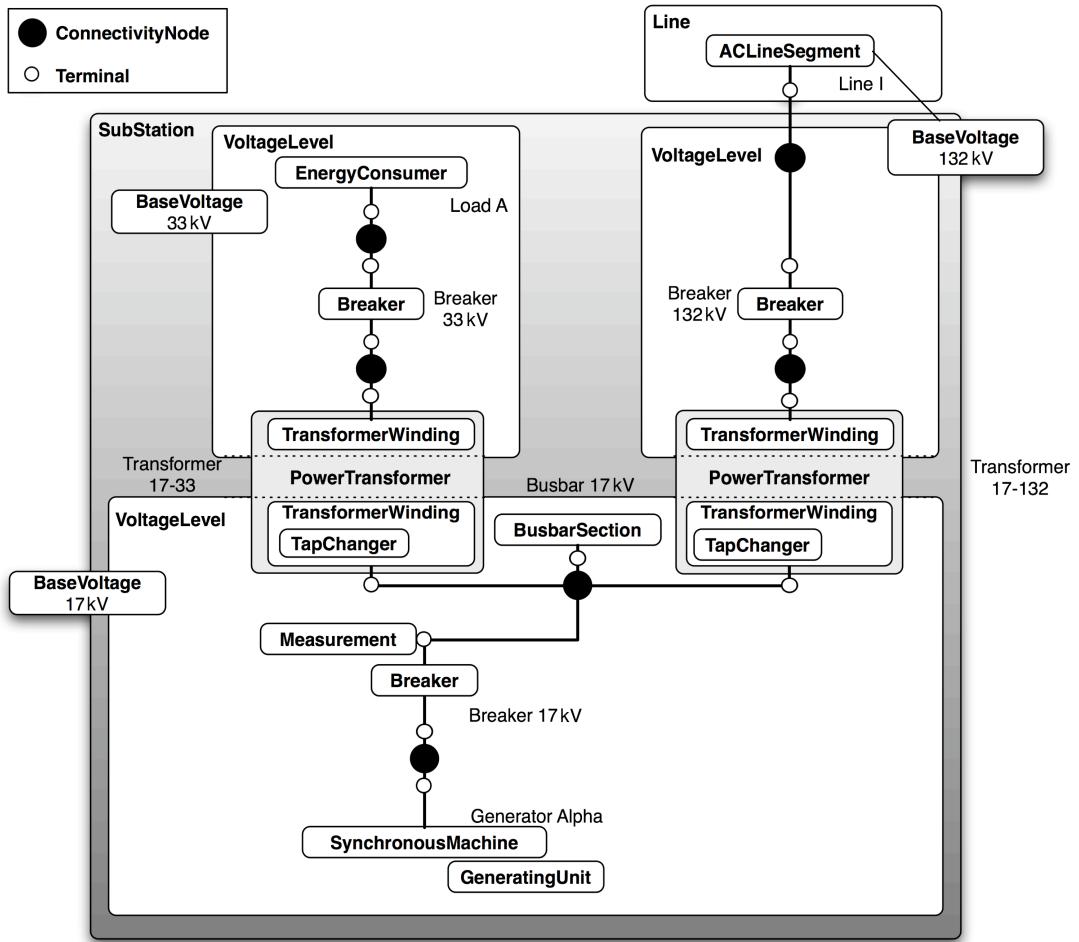


Figure 3.13 Example Circuit with full CIM Mappings

When fully converted to CIM objects, the original example circuit shown in Figure 3.9 is translated into the 45 CIM Objects shown in Figure 3.13. The BusbarSection's position may at first seem erroneous, but in the CIM the ConnectivityNodes are used to define the point of interconnection for pieces of equipment. As such, the BusbarSection object is used primarily to provide a point of association (via its terminal) for measurement objects measuring the voltage at that particular busbar in the system. This reflects the positioning of equipment in the physical system, since a voltage transformer will often measure voltages at the busbars within a substation.

This representation of the example network could be extended further with the addition of objects to represent control areas, equipment owners, measurement units and generation and load curves, but for now it is enough to understand how an existing network representation can be mapped to CIM objects.

3.4 IEC 61970-301 CIM Packages

As with any other complex class structure, classes in CIM are grouped together into packages dependent on their role within the power system. The core IEC 61970-301 standard contains eight main packages, and a global domain package used for defining

data types. The Core, Wires and Topology packages contain all the basic classes for defining the physical characteristics of a power network and, with the exception of the Measurement class, all the classes used in the CIM representation of the example circuit in Section 3.3 come from these three packages.

The Wires package defines classes that are required to represent the electrical components of a network, such as Transformers, Lines and Switches, while the Core and Topology packages define the interconnection of components: The Connectivity Node (contained in the Topology package) and Terminal (contained in the Core package).

These three packages alone, however, do not fully describe a functioning power system, but provide only the basic electrical characteristics of the equipment and describe how they are connected. To provide a detailed description of a network at an operational level, other classes are required to define the operation and additional characteristics of the equipment, both electrical and non-electrical.

The CIM is not only used for exchanging full power system models, as will be covered in more detail later on, the CIM is also used as a common model for defined business process messages. As such, a number of the packages contain classes that are used for business processes and not for defining the properties of a full power system model in CIM format.

3.4.1 Core

The Core package contains the parent PowerSystemResource class, from which all other classes concerned with the physical properties of the network inherit (including all classes relating to physical pieces of equipment, as well as *Equipment Containers* which are used for organising pieces of equipment into groups, such as specific VoltageLevels, or equipment contained within a specific Substation).

3.4.2 Wires

The Wires package defines all pieces of equipment electrically connected to the network, as well as supporting classes for defining additional properties and arrangement of objects. This includes classes for the components that are physically connected to the network at the points of power generation and consumption (Energy Consumer and Synchronous Machine, as previously mentioned), as well as several classes that detail the arrangement and settings for Power Transformers, properties for Lines (comprised of one or more Line Segments), and other pieces of conducting equipment including Switches, Busbar Sections and Regulating Conducting Equipment (Compensators).

3.4.3 Generation

The Generation package is split into two sub-packages, Production and GenerationDynamics.

Production

The Production package is used for defining various kinds of generators, and includes a class hierarchy for defining the components of Thermal and Hydro generators. The package also includes definitions of production costing information such as Cost Curves and Net to Gross curves. To define power generation unit in the CIM requires an

association of a production class object with a SynchronousMachine, a class contained within the Wires package.

GenerationDynamics

The GenerationDynamics package contains the description of Prime Movers, including turbine types, and classes that define various types of steam supplies, such as Pressurised or Boiling Water Reactors for nuclear power stations, and different types of Fossil Fuel Boilers for coal oil and gas fired boilers.

3.4.4 LoadModel

The LoadModel package deals with modelling energy consumers through curves and associated data. The EnergyConsumer class (and its subclasses) within the Wires package define the physical connection point between the network and customer. Instances of the EnergyConsumer class also contain associations to Load Demand Models and Schedules for non-conforming load (e.g. large industrial loads, or power station services).

3.4.5 Topology

The Topology package, together with the Terminal class, provides definitions of how equipment connects together in the form of Connectivity Nodes. The Topological Node class is comprised of Connectivity Nodes connected by closed switches (and for many applications can be considered analogous to a bus in a bus-branch representation). The Topological Island class contains all electrically connected Topological Nodes, and as such a fully interconnected power network should contain only one Topological Island.

3.4.6 Measurement

The Meas (Measurement) package is used to define the Measurements being taken from a particular Power System Resource. There are two ways of connecting Measurements:

The first option is to associate a measurement instance with a Power System Resource, which covers measurements not related to electrical connectivity including temperature or weight.

The second option, as described in Section 3.3.3 is to associate the Measurement with a Terminal. This is used for measurements dependent on connectivity, such as current or voltage where the Terminal defines the point of the network that the measurement is to be taken from. The Measurement class acts as a Current or Voltage Transformer for measuring the current or voltage at a point in the network, however it does not represent a piece of physical equipment.

3.4.7 Outage

The Outage package define schedules for the planned network configuration, and includes classes that associate with a Switch to define its state at a particular time. This is used primarily for defining business process messages, however it could be used to change network configurations at specific times during a simulation.

3.4.8 Protection

The Protection package defines the settings and parameters of pieces of protection equipment that operate Switches. The classes defined in this package are used to describe the behaviour of the Switch: its current limit, delay from detection of abnormal conditions to operation, maximum and minimum limits.

4 The eXtensible Markup Language (XML)

4.1 XML

XML, the eXtensible Markup Language, is a “universal format for structured documents and data”[3], which is quickly becoming the standard for storing machine-readable data in a structured, extensible format that is accessible over the internet. XML is actually a meta-language² that allows the user to design their own markup language to describe the structure of the data.

XML is a subset of SGML, the Standard Generalized Markup Language[6] designed for both on and offline storage and transferral of data. The data is encoded as plain text, thus allowing it to be both human and machine-readable and the use of standard encoding schemes makes it platform independent.

The XML syntax uses tags to denote the elements within the document. Each element is either expressed as an open and close tag containing data of the form:

```
<tag>...Contained Data...</tag>
```

Or with a single empty entry closed with a slash at the very end:

```
<tag/>
```

An entry may also contain its own attributes which are expressed in the form:

```
<tag attributeOne="something" attributeTwo="somethingElse"/> or
```

```
<tag attributeOne="something" attributeTwo="somethingElse">...</tag>
```

When an element has a start and end tag, any other elements contained within these two tags are classed as “children” of the parent element.

4.1.1 Simple XML Example

As an example, a simple XML tag-syntax to store a book can be created. The contents and properties of the book can then be expressed as XML, using self-descriptive tags of the form:

```
<book title="Introduction to XML" author="Alan McMorran">
  <revision number="2">
    <year>2006</year>
    <month>January</month>
    <day>1</day>
  </revision>
  <chapter title="Preface">
    <paragraph>Welcome to <italic>this</italic> book...</paragraph>
    <paragraph>...</paragraph>
    ...
    <paragraph>...and we shall continue</paragraph>
  </chapter>
  <chapter title="Introduction">
```

² A meta-language is a language used to describe a language (whether it be another language or itself).

```
<paragraph>To understand the uses...</paragraph>
...
</chapter>
</book>
```

Here the *book* element contains its own attribute to describe the *title* and *author*, with a child element to describe the *revision* of the book, plus several *chapter* elements. The *chapters* in turn contain elements for each *paragraph*, which themselves contain mixed data of other elements and text. Although to anybody with knowledge of the English language, the names of these tags make their semantics clear, the tag syntax and semantics must still be clearly defined if the data is to be interpreted correctly by an application.

4.1.2 XML Schema

While XML itself has no set tag-syntax or semantics, schemas can be defined for expressing almost any kind of data using XML notation. An application interpreting XML data must be given this knowledge of the syntax and semantics used, otherwise it will have trouble interpreting it. This requires the tag-syntax and semantics of the XML to be expressed as a schema, which provides constraints on the structure and contents of an XML document.

The most common formats for describing these schemas are in Document Type Definition (DTD)[4] format or the newer XML Schema[5]. The XML Schema defines the elements and attributes that can appear in a document; which elements are child elements; the number of allowed child elements for each element type; whether an element can include text (i.e. is an empty element or within an open and close tags); the data types for elements and attributes; whether their values are fixed; and if they have default values.

Using the previous book example, a simple XML Schema can be created to describe the elements within the document and the restrictions placed on them. This example schema is shown, along with additional comments, in Figure 4.1

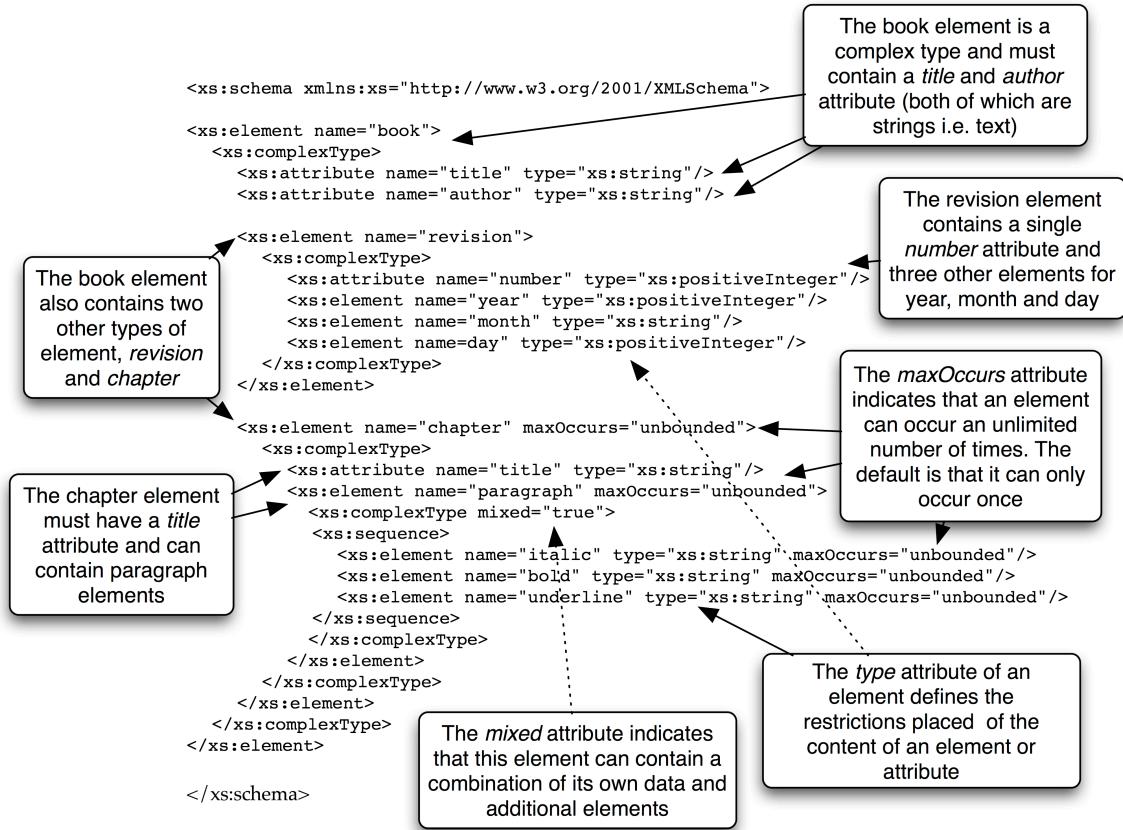


Figure 4.1 Annotated simple XML Schema Example describing the data within a book

The other notable feature of this document is the introduction of namespaces. In the example above, every element is prefixed by *xs*:. The document's root node contains an *xmlns:xs="http://www.w3.org/2001/XMLSchema"* attribute which indicates that every element prefixed with *xs* is an XML element that is part of the namespace identified by the Unique Resource Identifier (URI) *http://www.w3.org/2001/XMLSchema*³. An XML document can contain elements from multiple namespaces simultaneously, each of which denote a separate XML Schema with its own set of restrictions. For the previous example, the root node could become:

```
<xs:schema xmlns:xs="http://www.w3.org/2001/XMLSchema"
  xmlns:ab="http://www.other.com/2005/ABSchema"
  xmlns:yz="http://www.something.com/2004/YZ-Schema">
```

Indicating that the XML document may contain elements from the *http://www.other.com/2005/ABSchema* namespace, identified by an *ab* prefix, along with elements from *http://www.something.com/2004/YZ-Schema* namespace, identified by a *yz* prefix in addition to the original *http://www.w3.org/2001/XMLSchema* namespace.

³ The W3C is the World Wide Web Consortium, the governing body for web standards. Their domain is w3.org and as such W3C standards such as XML Schema and RDF use this domain as part of their unique resource identifier.

4.2 RDF

With a basic XML document there is no way to denote a link between two elements that are not parent or child. For instance, consider a library system containing entries for multiple books with information on their shelf position in the form:

```
<library name="Glasgow Library">
  <book title="History of Glasgow, 1900-1950" author="Walter Hannah">
    <position section="A" shelf="2"/>
  </book>
  <book title="A Brief History of Time" author="Stephen Hawking">
    <position section="E" shelf="4"/>
  </book>
  <book title="History of Glasgow, 1950-2000" author="Walter Hannah">
    <position section="A" shelf="2"/>
  </book>
</library>
```

Each *book* element is contained within the library as an independent entry, but should the user wish to add a link between the *History of Glasgow, 1900-1950* and *History of Glasgow, 1950-2000* books to indicate that reader may wish to read the former book before the latter, there is no standard way to do this using the basic XML constructs.

The Resource Document Framework (RDF)[9] is an XML schema used to provide a framework for data in an XML format by allowing relationships to be defined between XML nodes. Each element can be assigned a unique *ID* attribute under the RDF namespace <http://www.w3.org/1999/02/22-rdf-syntax-ns#> (which uses the *rdf* prefix). Adding a *resource* attribute to an element allows references to be made between elements by having its value refer to another element's *ID*.

4.2.1 Simple RDF Example

For the library example above, assigning an *ID* under the RDF namespace to each book allows the addition of *sequel* and *sequelTo* elements. These elements contain only a single resource attribute that point to another element within the document by referencing their *ID*.

To distinguish between the library elements and attributes, themselves governed by an XML Schema, and the RDF elements and attributes, an additional namespace <http://www.strath.ac.uk/libraries/2006/library-schema#> is added with the prefix *lib*. An RDF root element is also added with *xmlns* attributes to denote the namespaces and prefixes. The new Library RDF XML representation is shown below:

```
<rdf:RDF xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
  xmlns:lib="http://www.strath.ac.uk/libraries/2006/library-schema#">
  <lib:library lib:name="Glasgow Library">
    <lib:book lib:title="History of Glasgow, 1900-1950" lib:author="Walter
      Hannah" rdf:ID="_entry0001">
      <lib:position lib:section="A" lib:shelf="2"/>
      <lib:sequel rdf:resource="#_entry0003"/>
    </lib:book>
    <lib:book lib:title="A Brief History of Time" lib:author="Stephen Hawking"
      rdf:ID="_entry0002">
```

```

<lib:position lib:section="E" lib:shelf="4"/>
</book>
<lib:book lib:title="History of Glasgow, 1950-2000" lib:author="Walter
Hannah" rdf:ID="_entry0003">
    <lib:position lib:section="A" lib:shelf="2"/>
    <lib:sequelTo rdf:resource="#_entry0001"/>
</lib:book>
</lib:library>
</rdf:RDF>
```

As shown, the RDF provides a means of showing relationships between elements outwith the standard parent-child relationship. The schema contains additional elements that go beyond the simple *ID* and *resource* attribute as will be shown in next section, but it is these features of the framework that are utilised when expressing the CIM in XML format.

4.2.2 RDF Schema

While RDF provides a means of expressing simple statements about the relationship between resources, it does not define the vocabulary of these statements. The RDF Vocabulary Description Language, known as RDF Schema[10] provides the user with a means of describing specific kinds of resources or classes. The RDF Schema does not provide a vocabulary for a specific application's classes like *lib:sequel* or *lib:sequelTo*, or properties like *lib:title* and *lib:author*. Instead, the RDF Schema allows the user to describe these classes and properties themselves and indicate when they should be used together. For example, they may state that the property *lib:title* will be used in describing a *lib:book*, or that *lib:sequel* is an element of *lib:book* and should indicate a reference to another *lib:book* entry.

In essence, the RDF Schema provides a type system for RDF. The RDF Schema type system is similar to that of object-oriented programming languages such as Java, .NET and C++. Amongst other things, RDF Schema allows resources to be defined as instances of one or more classes and for these classes to be organised in a hierarchy.

For the previous example, the RDF Schema would, amongst others, contain entries to describe the class *book* and the properties *sequel* and *sequelTo*.

```

<rdfs:Class rdf:ID="book">
    <rdfs:label xml:lang="en">Book</rdfs:label>
    <rdfs:comment>A book contained within a library</rdfs:comment>
</rdfs:Class>

<rdf:Property rdf:ID="sequel">
    <rdfs:label xml:lang="en">Sequel</rdfs:label>
    <rdfs:comment>Indicates that the book has a sequel that is also within the
library</rdfs:comment>
    <rdfs:domain rdf:resource="#book"/>
    <rdfs:range rdf:resource="#book"/>
</rdf:Property>

<rdf:Property rdf:ID="sequelTo">
    <rdfs:label xml:lang="en">SequelTo</rdfs:label>
```

```

<rdfs:comment>Indicates that the book is the sequel to another book also
within the library</rdfs:comment>
<rdfs:domain rdf:resource="#book"/>
<rdfs:range rdf:resource="#book"/>
</rdf:Property>

```

Here, the class of *book* is defined, then the two properties *sequel* and *sequelTo* are defined. Each of these properties has their domain (the class the property is within) referencing the *book* class, as does their range (the class of element the property refers to). Should the library schema be extended so that instead of just having a *book* element, fictional novels could be differentiated with a separate *novel* element that, when modelled in UML, would be a simple sub-class of the existing *book* class. This can be represented in RDF Schema as:

```

<rdfs:Class rdf:ID="novel">
  <rdfs:label xml:lang="en">Novel</rdfs:label>
  <rdfs:comment>A fictional book</rdfs:comment>
  <rdfs:subClassOf rdf:resource="#book"/>
</rdfs:Class>

```

The RDF, combined with RDF Schema provides a mechanism for expressing a basic class hierarchy as an XML schema by specifying the basic relationship between classes and properties. This then allows a set of objects to be expressed as XML using a defined schema that retain their relationships and class hierarchy.

4.3 CIM RDF XML

Since the RDF and RDF Schema provide a means of mapping an object-oriented design to XML the CIM class structure can be mapped in a similar way. Existing tools can automatically generate an RDF Schema from the original CIM UML model. Using the previous example of the Transformer class hierarchy shown in Figure 3.11, the resulting RDF Schema takes the form:

```

<rdfs:Class rdf:ID="PowerSystemResource">
  <rdfs:label xml:lang="en">PowerSystemResource</rdfs:label>
  <rdfs:subClassOf rdf:resource="#Naming"/>
</rdfs:Class>

<rdfs:Class rdf:ID="Equipment">
  <rdfs:label xml:lang="en">Equipment</rdfs:label>
  <rdfs:subClassOf rdf:resource="#PowerSystemResource"/>
</rdfs:Class>

<rdfs:Class rdf:ID="ConductingEquipment">
  <rdfs:label xml:lang="en">ConductingEquipment</rdfs:label>
  <rdfs:subClassOf rdf:resource="#Equipment"/>
</rdfs:Class>

<rdfs:Class rdf:ID="PowerTransformer">
  <rdfs:label xml:lang="en">PowerTransformer</rdfs:label>
  <rdfs:subClassOf rdf:resource="#Equipment"/>
</rdfs:Class>

```

```

<rdfs:Class rdf:ID="TransformerWinding">
  <rdfs:label xml:lang="en">TransformerWinding</rdfs:label>
  <rdfs:subClassOf rdf:resource="#ConductingEquipment"/>
</rdfs:Class>

<rdfs:Class rdf:ID="TapChanger">
  <rdfs:label xml:lang="en">TapChanger</rdfs:label>
  <rdfs:subClassOf rdf:resource="#PowerSystemResource"/></rdfs:Class>
</rdfs:Class>

<rdf:Property rdf:ID="TransformerWinding.MemberOf_PowerTransformer">
  <rdfs:label xml:lang="en">MemberOf_PowerTransformer</rdfs:label>
  <rdfs:domain rdf:resource="#TransformerWinding"/>
  <rdfs:range rdf:resource="#PowerTransformer"/>
</rdf:Property>

<rdf:Property rdf:ID="TapChanger.TransformerWinding">
  <rdfs:label xml:lang="en">TransformerWinding</rdfs:label>
  <rdfs:domain rdf:resource="#TapChanger"/>
  <rdfs:range rdf:resource="#TransformerWinding"/>
</rdf:Property>

```

Each CIM Class has a corresponding *rdfs:Class* entry, while the two aggregation relationships are expressed as RDF *Property* elements with the appropriate domains and ranges. The entire CIM Class structure can be expressed in this manner, and then this RDF Schema can be used to express a CIM power system model as RDF XML.

4.3.1 CIM RDF XML Example

As an example, the simple Transformer example from Figure 3.12 will be extended to include attributes for each object. This produces four objects with their own internal data as shown in Figure 4.2 below:

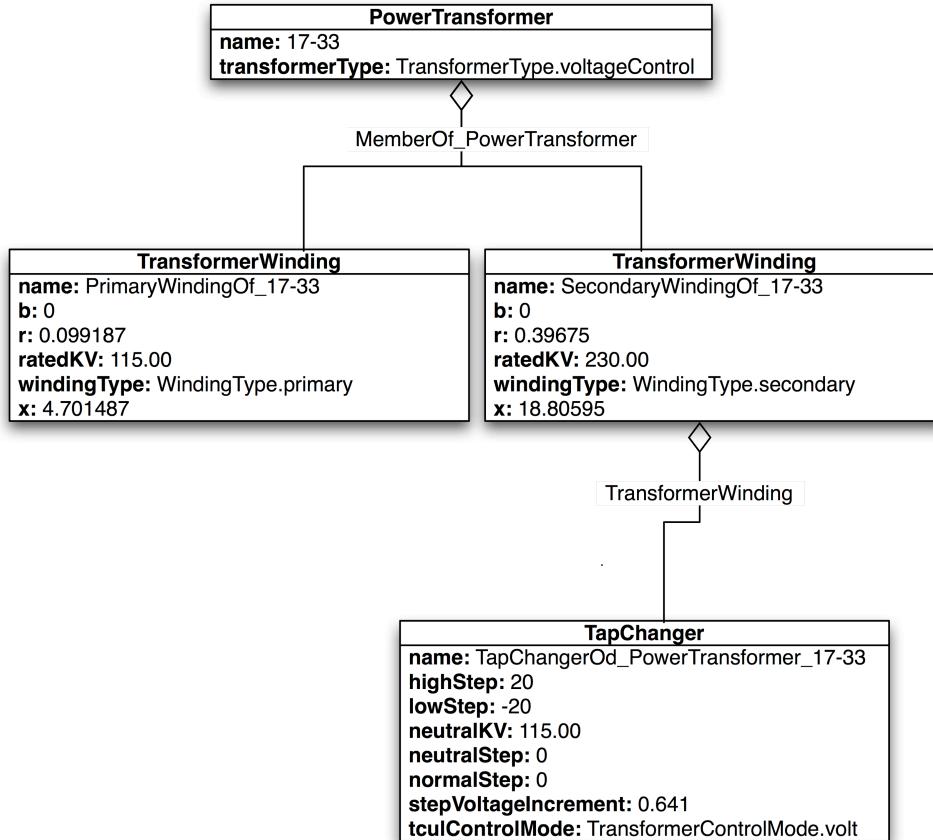


Figure 4.2 Transformer shown as four CIM Objects with attributes

Each of these objects can then be expressed as an XML node using the CIM RDF Schema given the namespace <http://iec.ch/TC57/2003/CIM-schema-cim10#> and prefix `cim`:

```

<rdf:RDF xmlns:cim="http://iec.ch/TC57/2003/CIM-schema-cim10#"
  xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#">

```

```

<cim:PowerTransformer rdf:ID="PowerTransformer_1733">
  <cim:PowerTransformer.transformerType
  rdf:resource="http://iec.ch/TC57/2003/CIM-schema-
  cim10#TransformerType.voltageControl"/>
  <cim:Naming.name>17-33</cim:Naming.name>
</cim:PowerTransformer>

<cim:TransformerWinding rdf:ID="PrimaryWindingOf_PowerTransformer_1733">
  <cim:TransformerWinding.b>0</cim:TransformerWinding.b>
  <cim:TransformerWinding.r>0.099187</cim:TransformerWinding.r>
  <cim:TransformerWinding.ratedKV>115.00</cim:TransformerWinding.ratedKV>
  <cim:TransformerWinding.windingType
  rdf:resource="http://iec.ch/TC57/2003/CIM-schema-cim10#WindingType.primary"/>
  <cim:TransformerWinding.x>4.701487</cim:TransformerWinding.x>
  <cim:TransformerWinding.MemberOf_PowerTransformer
  rdf:resource="#PowerTransformer_302"/>
  <cim:Naming.name>PrimaryWindingOf_17-33</cim:Naming.name>
</cim:TransformerWinding>

```

```

<cim:TransformerWinding rdf:ID="SecondaryWindingOf_PowerTransformer_1733">
  <cim:TransformerWinding.b>0</cim:TransformerWinding.b>
  <cim:TransformerWinding.r>0.39675</cim:TransformerWinding.r>
  <cim:TransformerWinding.ratedKV>230.00</cim:TransformerWinding.ratedKV>
  <cim:TransformerWinding.windingType
  rdf:resource="http://iec.ch/TC57/2003/CIM-schema-
  cim10#WindingType.secondary"/>
  <cim:TransformerWinding.x>18.80595</cim:TransformerWinding.x>
  <cim:TransformerWinding.MemberOf_PowerTransformer
  rdf:resource="#PowerTransformer_302"/>
  <cim:Naming.name>SecondaryWindingOf_17-33</cim:Naming.name>
</cim:TransformerWinding>

<cim:TapChanger rdf:ID="TapChangerOf_PowerTransformer_1733">
  <cim:TapChanger.highStep>20</cim:TapChanger.highStep>
  <cim:TapChanger.lowStep>-20</cim:TapChanger.lowStep>
  <cim:TapChanger.neutralKV>115.00</cim:TapChanger.neutralKV>
  <cim:TapChanger.neutralStep>0</cim:TapChanger.neutralStep>
  <cim:TapChanger.normalStep>0</cim:TapChanger.normalStep>
  <cim:TapChanger.stepVoltageIncrement>0.641</cim:TapChanger.stepVoltageIncre
  ment>
  <cim:TapChanger.tculControlMode
  rdf:resource="http://iec.ch/TC57/2003/CIM-
  schema-cim10#TransformerControlMode.volt"/>
  <cim:TapChanger.TransformerWinding
  rdf:resource="#PrimaryWindingOf_PowerTransformer_302"/>
  <cim:Naming.name>TapChangerOf_PowerTransformer_17-33</cim:Naming.name>
</cim:TapChanger>

</rdf:RDF>

```

The PowerTransformer.transformerType and TapChanger.tculControlMode elements do not refer to other nodes within the document; instead their values are of an enumerated type. Enumerated types consist of a fixed set of legal values (e.g. for a variable of type Days, the enumerated type would be: Sunday, Monday, Tuesday, Wednesday, Thursday, Friday, Saturday and any variable of this type must have one of these values). Within the CIM there are certain class attributes that are also enumerated types and do not contain a node value but instead refer to an enumerated type within its RDF Schema.

This combination of the CIM, XML, RDF and RDF Schema allows an entire CIM power system model to be expressed in a standard, cross-platform plain-text format that is both human and machine readable and extensible.

5 XML Messaging

As well as exchanging full power system model data as CIM RDF XML, the other main application of the CIM is as a common semantic model for enterprise application integration.

5.1 Existing Inter-Application Communication Infrastructure

Within large companies there will be a number of computer applications that must communicate with each other. This often results in a large number of point-to-point links using custom formats and protocols to exchange data between software applications from a number of vendors. Adding a new application to the system requires additional communication links to be defined and implemented, further increasing the complexity of the overall system with a corresponding financial penalty.

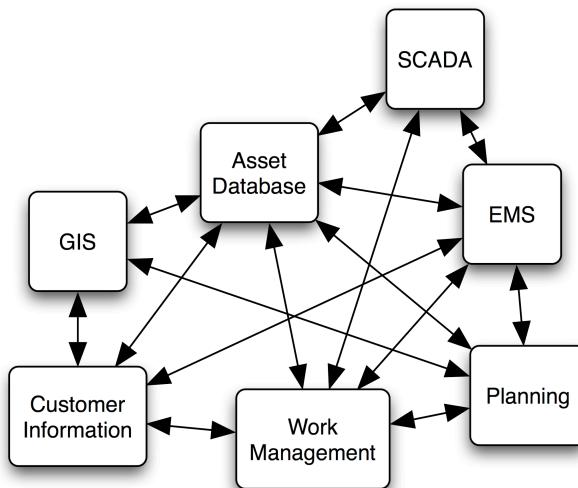


Figure 5.1 Communication links between enterprise applications

As illustrated in Figure 5.1, even for a small section of the overall IT system, this can result in a large number of inter-application communication links. As companies expand their IT infrastructure or replace existing applications with products from other vendors they must define new interfaces for each communication link, a process that is both time consuming and expensive.

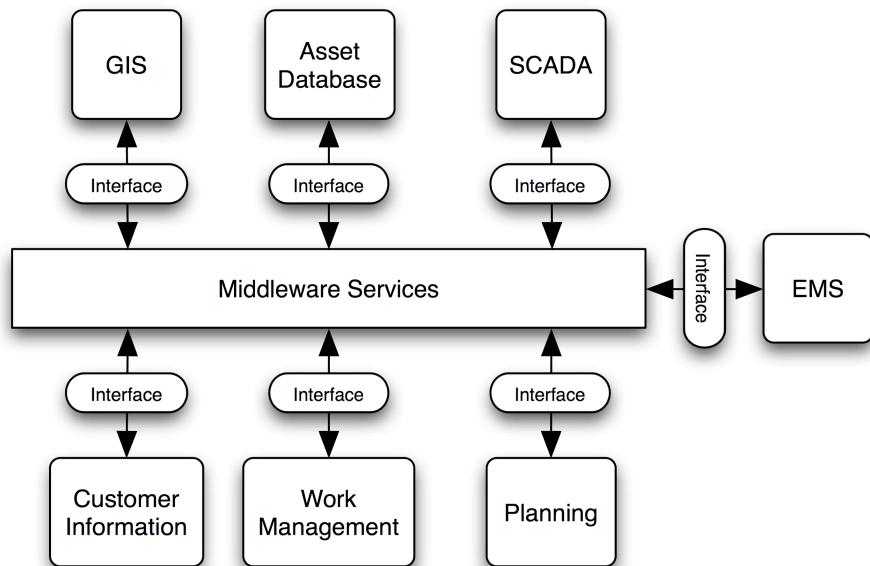


Figure 5.2 Enterprise Application Bus model for inter-application communication

5.2 The Message Bus Concept

Enterprise Application Integration (EAI)[7] replaces these dedicated links with a single communication link called a “message bus”. Using middleware services, this provides a mechanism for applications to communicate using a pre-defined message format and requires only a single interface to be written for each application.

The CIM provides the common semantic model used to construct the messages that are used for communication between the applications[8]. This requires each application to map its external interface to the CIM class structure allowing the inter-application messages to be defined in the CIM.

These messages, in XML format, use a restricted CIM XML Schema to define the payloads of the messages. This takes the standard CIM Schema, itself created from the CIM UML class structure and restricts the multiplicity of associations and required attributes.

5.3 Mapping Application Interfaces to the CIM

This can be illustrated using a simple example. An EMS application’s external interface requires the user to access data on the transformers within the system. The EMS application’s interface attributes are:

- *TRANS_NAME* – The Transformer’s name
- *WINDINGA_R* – The Transformer’s primary winding resistance
- *WINDINGA_X* – The Transformer’s primary winding reactance
- *WINDINGB_R* – The Transformer’s secondary winding resistance
- *WINDINGB_X* – The Transformer’s secondary winding reactance
- *WINDINGA_V* – The Transformer’s primary winding voltage
- *WINDINGB_V* – The Transformer’s secondary winding voltage

Each of these attributes can be mapped to a corresponding attribute within a CIM class, resulting in an interface to CIM mapping.

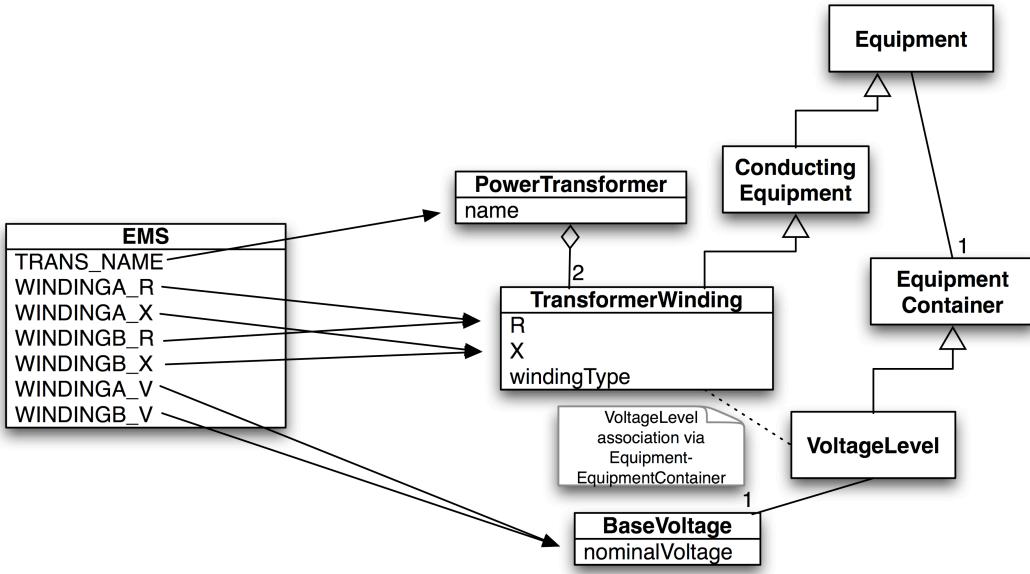


Figure 5.3 CIM Interface Mapping

This mapping, shown in Figure 5.3, highlights that although the two windings have separate names in the interface, they map to the same attributes within the CIM class structure. The aggregation relationship between the **PowerTransformer** and **TransformerWinding** class has, however, been changed from a 0..n multiplicity to 2 (since in this example the EMS represents all transformers as having two windings). This means that there must be two instances of the **TransformerWinding** class present in the message, with the **windingType** attribute then used to differentiate between the primary and secondary windings.

The voltage for each winding is contained with the *nominalVoltage* attribute of the **BaseVoltage** class. The **BaseVoltage** instance is associated with the **TransformerWinding** via the **VoltageLevel** class whose relationship with **TransformerWinding** is defined in the **Equipment** - **EquipmentContainer** association further up the class hierarchy. This is because the **TransformerWinding** class itself does not contain a direct **VoltageLevel** association within the CIM, but instead inherits a **MemberOf_Equipment** container association from the **Equipment** class (via **ConductingEquipment**), and since **VoltageLevel** is a subclass of **EquipmentContainer** this can be used to provide the required association to **VoltageLevel**.

Both the **EquipmentContainer** and **BaseVoltage** associations have their multiplicity changed from 0..1 to 1 requiring each **VoltageLevel** to have one **BaseVoltage** instance associated with it and each piece of **Equipment** (in this case **TransformerWinding**) to have an association to a single **EquipmentContainer** instance.

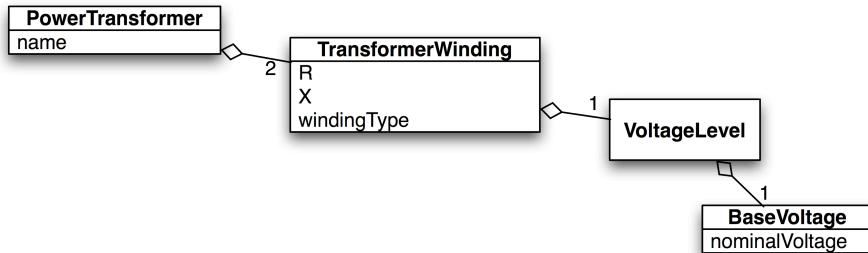


Figure 5.4 Message Payload as UML

5.4 Constructing a Message Payload

This message payload can be further restricted by changing the associations to aggregations and removing the parent classes since they are not required by the actual message content. Thus the message payload can be represented as the modified class structure shown in Figure 5.4.

Thus a single two winding transformer containing the desired attributes can be represented in XML as:

```
<cim:PowerTransformer>
  <cim:Naming.name>Transformer SGT1</cim:Naming.name>
  <cim:PowerTransformer.Contains_TransformerWindings>
    <cim:TransformerWinding.r>0.23</cim:TransformerWinding.r>
    <cim:TransformerWinding.x>0.78</cim:TransformerWinding.x>
    <cim:TransformerWinding.windingType>WindingType.primary
      </cim:TransformerWinding.windingType>
    <cim:Equipment.MemberOf_EquipmentContainer>
      <cim:VoltageLevel.BaseVoltage>
        <cim:BaseVoltage.nominaVoltage>400
          </cim:BaseVoltage.nominalVoltage>
        </cim:VoltageLevel.BaseVoltage>
      </cim:Equipment.MemberOf_EquipmenContainer>
    </cim:PowerTransformer.Contains_TransformerWindings>
    <cim:PowerTransformer.Contains_TransformerWindings>
      <cim:TransformerWinding.r>0.46</cim:TransformerWinding.r>
      <cim:TransformerWinding.x>0.87</cim:TransformerWinding.x>
      <cim:TransformerWinding.windingType>WindingType.secondary
        </cim:TransformerWinding.windingType>
      <cim:Equipment.MemberOf_EquipmentContainer>
        <cim:VoltageLevel.BaseVoltage>
          <cim:BaseVoltage.nominaVoltage>275
            </cim:BaseVoltage.nominalVoltage>
          </cim:VoltageLevel.BaseVoltage>
        </cim:Equipment.MemberOf_EquipmenContainer>
      </cim:PowerTransformer.Contains_TransformerWindings>
    </cim:PowerTransformer>
```

This XML message in turn has an XML Schema to describe the payload contents:

```

<xs:schema xmlns:cim="cimBase" xmlns:xs="http://www.w3.org/2001/XMLSchema">

<xs:element minOccurs="1" maxOccurs="1" name="PowerTransformer">
  <xs:complexType>
    <xs:complexContent>
      <xs:extension base="cim:PowerTransformer">
        <xs:sequence>
          <xs:element minOccurs="1" maxOccurs="1"
            name="Naming.name" type="xs:string"/>
          <xs:element minOccurs="2" maxOccurs="2"
            name="PowerTransformer.Contains_TransformerWindings">
            <xs:complexType>
              <xs:complexContent>
                <xs:extension base="cim:TransformerWinding">
                  <xs:sequence>
                    <xs:element minOccurs="1" maxOccurs="1"
                      name="TransformerWinding.r" type="xs:float"/>
                    <xs:element minOccurs="1" maxOccurs="1"
                      name="TransformerWinding.x" type="xs:float"/>
                    <xs:element minOccurs="1" maxOccurs="1"
                      name="TransformerWinding.windingType" type="cim:WindingType"/>
                    <xs:element minOccurs="1" maxOccurs="1"
                      name="TransformerWinding.MemberOf_EquipmentContainer">
                      <xs:complexType>
                        <xs:complexContent>
                          <xs:extension base="cim:VoltageLevel">
                            <xs:sequence>
                              <xs:element minOccurs="1" maxOccurs="1"
                                name="VoltageLevel.BaseVoltage">
                                <xs:complexType>
                                  <xs:complexContent>
                                    <xs:extension base="cim:BaseVoltage">
                                      <xs:sequence>
                                        <xs:element minOccurs="1" maxOccurs="1"
                                          name="BaseVoltage.nominalVoltage" type="xs:float"/>
                                      </xs:sequence>
                                    </xs:extension>
                                  </xs:complexContent>
                                </xs:extension>
                              </xs:sequence>
                            </xs:extension>
                          </xs:complexContent>
                        </xs:extension>
                      </xs:complexType>
                    </xs:element>
                  </xs:sequence>
                </xs:extension>
              </xs:complexContent>
            </xs:extension>
          </xs:element>
        </xs:sequence>
      </xs:extension>
    </xs:complexContent>
  </xs:complexType>
</xs:element>
</xs:schema>

```

This schema refers to another *cimBase* schema that contains the definitions for the CIM classes, their attributes and associations and for the enumerations and data types such as *WindingType*.

The major difference between the CIM XML Message and the CIM RDF XML is that instead of having each CIM object as an independent XML element that is then linked using the RDF *ID* and *resource* attribute, the elements are contained within each other. This way a PowerTransformer element contains child elements to describe the TransformerWindings, that in turn contain child elements to denote the voltage.

For the scope of the message, this means that the PowerTransformer element contains all the required data making it simpler to then transform the element into another format if that is required. Importing and converting CIM RDF XML data can more challenging, but given the highly interconnected nature of that representation it is not possible to represent it as nested XML elements for anything other than the most simple network models.

5.5 XML Messaging Summary

This example has shown how a simple portion of an application's interface can be mapped to the CIM class structure and then used to construct a simple XML message payload. Real-world examples often use tens or even hundreds of elements to construct a message payload. The benefit of this approach is that when every application within the system is mapped to this common model it becomes far simpler for applications to communicate. The CIM provides not only a common data format but crucially provides a common semantic model, which provides consensus on the interpretation of each class and attribute.

This CIM XML messaging approach has been applied by a number of large utilities both in the UK and the USA and has proven to be a flexible and scalable system. As well as exchanging data on logical power system components, IEC 61968-11 contains extensions to the core 61970-301 CIM packages and classes to allow the definition of business processes such as work scheduling, customer invoicing and financial trading arrangement as a CIM XML messages.

Much of the work being undertaken to extend the CIM is concerned with using it to define message payloads for exchanging between applications in an EAI environment. This is why only a small subset of the overall 61970-301 and 61968-11 CIM class structure is used when representing a logical power systems model.

6 Summary

This paper has introduced the concept of classes and class hierarchies along with their basic relationships that define how classes relate to each other: inheritance, association, aggregation and composition. The benefits of using this approach to define the components of a power system were then demonstrated along with an example of how a simple power system, represented as a line diagram, can be mapped to CIM Objects. The extensible markup language, its resource document framework subset and schemas were then introduced to demonstrate how the CIM class structure is mapped and the data encapsulated in an XML format. Finally the primary uses of the CIM were discussed: for encapsulating entire power system models as CIM RDF XML; and exchanging data between applications as CIM XML Messages.

7 References

- [1] "IEC 61970 Energy management system application program interface (EMS-API) - Part 301: Common Information Model (CIM) Base", IEC, Edition 1.0, November 2003
- [2] "IEC 61968 Application integration at electric utilities - System interfaces for distribution management - Part 11: Common Information Model (CIM)", IEC Draft
- [3] W3C Recommendation, "Extensible Markup Language" Version 1.0, October 2000, available at <http://www.w3.org/TR/REC-xml>
- [4] W3C Recommendation, "Extensible Markup Language: Prolog and Document Type Declaration" Version 1.0, October 2000, available at <http://www.w3.org/TR/REC-xml/#sec-prolog-dtd>
- [5] W3C Recommendation, "XML Schema Part 0: Primer Second Edition", October 2004, available at <http://www.w3.org/TR/xmlschema-0/>
- [6] Information processing - Text and office systems - Standard Generalized Markup Language (SGML), ISO 8879
- [7] D. Linthicum, Enterprise Application Integration, Addison Wesley Longman, Reading, Massachusetts, 2000.
- [8] G. Robinson, "Model Driven Integration (MDI) for Electric Utilities", Proceedings of Distributech, Miami Beach, Florida, USA, March 2002
- [9] W3C Recommendation, "Resource Description Framework", February 1999, available at <http://www.w3.org/TR/REC-rdf-syntax/>
- [10] W3C Recommendation, "RDF Vocabulary Description Language 1.0: RDF Schema", Version 1.0, February 2004, available at <http://www.w3.org/TR/rdf-schema/>
- [11] Arnold deVos, "Simplified RDF Syntax for Power System Model Exchange", Langdale Consultants, available at <http://www.langdale.com.au/DAF/PSModelExchange.pdf>